

```
In [1]: !python --version
```

```
Python 3.11.5
```

```
In [2]: from __future__ import print_function, division

import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
from torch.autograd import Variable
import numpy as np
import torchvision
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
import cv2
import time
import os
import copy
import skimage
from tensorflow.keras.utils import to_categorical

imageSize=200
train_dir = "C:/Users/sumed/Downloads/archive (12)/OCT2017/train/"
test_dir = "C:/Users/sumed/Downloads/archive (12)/OCT2017/test/"
val_dir = "C:/Users/sumed/Downloads/archive (12)/OCT2017/val/"
# ['DME', 'CNV', 'NORMAL', '.DS_Store', 'DRUSEN']
from tqdm import tqdm
```


```
In [3]: from pathlib import Path
normal_dir = Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/NORMAL")

# Gather all image file paths and select the first 20
image_paths = list(normal_dir.glob('*'))[:20]

# Set up the plotting parameters
plt.figure(figsize=(10, 8)) # Adjust size as needed
for i, image_path in enumerate(image_paths):
    # Read and resize each image
    image = cv2.imread(str(image_path))
    image = cv2.resize(image, (128, 128))
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert color format

    # Plot each image in a 4x5 grid
    plt.subplot(4, 5, i + 1)
    plt.imshow(image)
    plt.axis('off') # Hide axes for a cleaner display

plt.tight_layout() # Adjust layout to fit images well
plt.show()
```




```
In [4]: cnv_dir = Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/CNV")

# Gather all image file paths and select the first 20
image_paths_cnv = list(cnv_dir.glob('*'))[:20]

# Set up the plotting parameters
plt.figure(figsize=(10, 8)) # Adjust size as needed
for i, image_path in enumerate(image_paths_cnv):
    # Read and resize each image
    image_cnv = cv2.imread(str(image_path)) # Use the correct image path here
    if image_cnv is None:
        print(f"Failed to load image: {image_path}")
        continue # Skip this image if it can't be loaded
    image_cnv = cv2.resize(image_cnv, (128, 128)) # Resize the image
    image_cnv = cv2.cvtColor(image_cnv, cv2.COLOR_BGR2RGB) # Convert color format

    # Plot each image in a 4x5 grid
    plt.subplot(4, 5, i + 1)
    plt.imshow(image_cnv)
    plt.axis('off') # Hide axes for a cleaner display

plt.tight_layout() # Adjust layout to fit images well
plt.show()
```



```
In [5]: dme_dir = Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/DME")

# Gather all image file paths and select the first 20
image_paths_dme = list(dme_dir.glob('*'))[:20]


# Set up the plotting parameters
plt.figure(figsize=(10, 8)) # Adjust size as needed
for i, image_path in enumerate(image_paths_dme):
    # Read and resize each image
    image_dme = cv2.imread(str(image_path))

    if image_dme is None:
        print(f"Failed to load image: {image_path}")
        continue # Skip this image if it can't be loaded

    image_dme = cv2.resize(image_dme, (128, 128)) # Resize the image
    image_dme = cv2.cvtColor(image_dme, cv2.COLOR_BGR2RGB) # Convert color format

    # Plot each image in a 4x5 grid
    plt.subplot(4, 5, i + 1)
    plt.imshow(image_dme)
    plt.axis('off') # Hide axes for a cleaner display

plt.tight_layout() # Adjust layout to fit images well
plt.show()
```

```
In [6]: drusen_dir = Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/DRUSEN")

# Gather all image file paths and select the first 20
image_paths_drusen = list(drusen_dir.glob('*'))[:20]


# Set up the plotting parameters
plt.figure(figsize=(10, 8)) # Adjust size as needed
for i, image_path in enumerate(image_paths_drusen):
    # Read and resize each image
    image_drusen = cv2.imread(str(image_path))

    if image_drusen is None:
        print(f"Failed to load image: {image_path}")
        continue # Skip this image if it can't be loaded

    image_drusen = cv2.resize(image_drusen, (128, 128)) # Resize the image
    image_drusen = cv2.cvtColor(image_drusen, cv2.COLOR_BGR2RGB) # Convert color format

    # Plot each image in a 4x5 grid
    plt.subplot(4, 5, i + 1)
    plt.imshow(image_drusen)
    plt.axis('off') # Hide axes for a cleaner display

plt.tight_layout() # Adjust layout to fit images well
plt.show()
```



```
In [7]: # Define the directories for each class
class_dirs = {
    'NORMAL': Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/NORMAL"),
    'DME': Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/DME"),
    'CNV': Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/CNV"),
    'DRUSEN': Path("C:/Users/sumed/Downloads/archive (12)/OCT2017/train/DRUSEN")
}


# Initialize a dictionary to store the counts
class_counts = {}

# Loop through each class directory and count the images
for class_name, class_dir in class_dirs.items():
    class_counts[class_name] = len(list(class_dir.glob('*'))) # Count files in each cla

# Plot the results in a bar graph
plt.figure(figsize=(8, 6))
colors = plt.cm.Purples(np.linspace(0.4, 0.8, len(class_counts)))

plt.bar(class_counts.keys(), class_counts.values(), color=colors)

# Add labels and title
plt.xlabel('Class')
plt.ylabel('Number of Images')
plt.title('Number of Images per Class')
plt.show()
```



The weights are computed to make the model focus more on the classes that are less frequent. Since the weights are close to 1, it suggests that your dataset does not have a severe class imbalance. The model will treat DME (which appears to have slightly more samples) with a slightly lower weight, and NORMAL, CNV, and DRUSEN will get roughly equal treatment with weights near 1.

```
In [8]: img_size = (128, 128)
        batch_size = 32
```

```
In [9]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
        train_datagen = ImageDataGenerator(rescale=1./255, rotation_range=10, width_shift_range=
        val_datagen = ImageDataGenerator(rescale=1./255)
        test_datagen = ImageDataGenerator(rescale=1./255)
```

```
In [10]: train_generator = train_datagen.flow_from_directory(train_dir, target_size=img_size, bat
        val_generator = val_datagen.flow_from_directory(val_dir, target_size=img_size, batch_siz
        test_generator = test_datagen.flow_from_directory(test_dir, target_size=img_size, batch_

Found 83484 images belonging to 4 classes.
Found 32 images belonging to 4 classes.
Found 968 images belonging to 4 classes.
```

```
In [11]: from sklearn.utils.class_weight import compute_class_weight

        # Extract class indices and their corresponding counts from the training generator
        class_indices = train_generator.class_indices
        classes = list(class_indices.keys()) # Class names
        y_train = train_generator.classes   # True class labels from the generator

        # Compute class weights
```

```
class_weights = compute_class_weight('balanced', classes=np.unique(y_train), y=y_train)

# Convert to dictionary format for easy mapping to class names
class_weights_dict = dict(zip(class_indices.values(), class_weights))

print(f"Class Weights: {class_weights_dict}")
```

Class Weights: {0: 0.5609729875016799, 1: 1.8391787099048291, 2: 2.4223537604456826, 3: 0.7931217936538096}

```
In [11]: import tensorflow as tf

# Build CNN model
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Conv2D(128, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(4, activation='softmax')
])
```

C:\Users\sumed\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [12]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [13]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_1 (Conv2D)	(None, 61, 61, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 128)	3,211,392
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 4)	516

Total params: 3,305,156 (12.61 MB)

Trainable params: 3,305,156 (12.61 MB)

Non-trainable params: 0 (0.00 B)

```
In [16]: model_training = model.fit(train_generator, validation_data=val_generator, epochs=10)
```

Epoch 1/10

```
/opt/conda/lib/python3.10/site-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
  self._warn_if_super_not_called()
```

2609/2609 ————— 1929s 737ms/step - accuracy: 0.6299 - loss: 0.9498 - val_accuracy: 0.5312 - val_loss: 1.0733

Epoch 2/10

2609/2609 ————— 1819s 696ms/step - accuracy: 0.7601 - loss: 0.6444 - val_accuracy: 0.8438 - val_loss: 0.5039

Epoch 3/10

2609/2609 ————— 1776s 680ms/step - accuracy: 0.8437 - loss: 0.4470 - val_accuracy: 0.8125 - val_loss: 0.5071

Epoch 4/10

2609/2609 ————— 1716s 657ms/step - accuracy: 0.8739 - loss: 0.3688 - val_accuracy: 0.9688 - val_loss: 0.0807

Epoch 5/10

2609/2609 ————— 1809s 675ms/step - accuracy: 0.8839 - loss: 0.3461 - val_accuracy: 0.9375 - val_loss: 0.1014

Epoch 6/10

2609/2609 ————— 1680s 643ms/step - accuracy: 0.8956 - loss: 0.3128 - val_accuracy: 1.0000 - val_loss: 0.0518

Epoch 7/10

2609/2609 ————— 1708s 645ms/step - accuracy: 0.9015 - loss: 0.2974 - val_accuracy: 1.0000 - val_loss: 0.0556

Epoch 8/10

2609/2609 ————— 1674s 641ms/step - accuracy: 0.9064 - loss: 0.2811 - val_accuracy: 0.9375 - val_loss: 0.0938

Epoch 9/10

2609/2609 ————— 1676s 642ms/step - accuracy: 0.9084 - loss: 0.2800 - val_accuracy: 0.9688 - val_loss: 0.0318

Epoch 10/10

2609/2609 ————— 1705s 643ms/step - accuracy: 0.9104 - loss: 0.2703 - val_accuracy: 0.9688 - val_loss: 0.0477

```
In [18]: test_loss, test_accuracy = model.evaluate(test_generator)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
```

31/31 ————— 10s 316ms/step - accuracy: 0.9712 - loss: 0.1164

Test Loss: 0.1068936362862587

Test Accuracy: 97.21%

```
In [19]: # Predicting on the test set
test_predictions = model.predict(test_generator)
test_predictions_labels = np.argmax(test_predictions, axis=1)
true_labels = test_generator.classes
```

31/31 ————— 7s 205ms/step

```
In [22]: import seaborn as sns

# Plot training and validation accuracy
sns.set_theme(style="whitegrid")

# Extract training history
history = model_training.history
```



```

# Create subplots for accuracy and loss
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

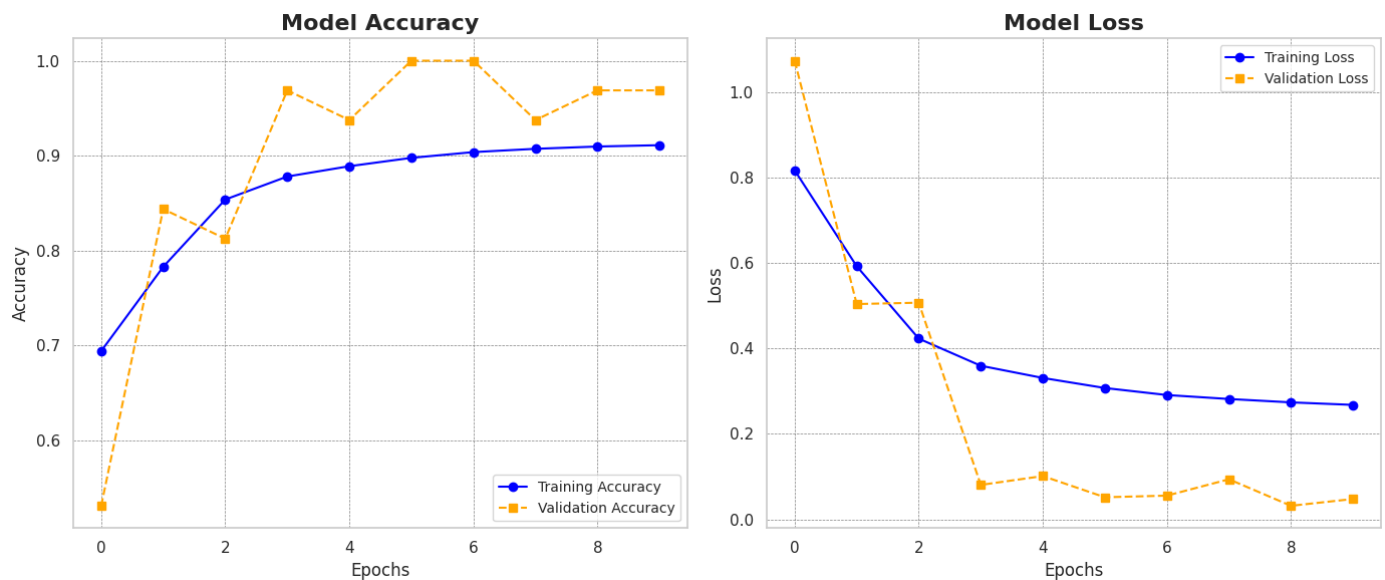
# Plot training and validation accuracy
axes[0].plot(history['accuracy'], label='Training Accuracy', marker='o', linestyle='-',
axes[0].plot(history['val_accuracy'], label='Validation Accuracy', marker='s', linestyle='--', c
axes[0].set_title('Model Accuracy', fontsize=16, weight='bold')
axes[0].set_xlabel('Epochs', fontsize=12)
axes[0].set_ylabel('Accuracy', fontsize=12)
axes[0].legend(loc='lower right', fontsize=10)
axes[0].grid(color='gray', linestyle='--', linewidth=0.5)

# Plot training and validation loss
axes[1].plot(history['loss'], label='Training Loss', marker='o', linestyle='-', color='b'
axes[1].plot(history['val_loss'], label='Validation Loss', marker='s', linestyle='--', c
axes[1].set_title('Model Loss', fontsize=16, weight='bold')
axes[1].set_xlabel('Epochs', fontsize=12)
axes[1].set_ylabel('Loss', fontsize=12)
axes[1].legend(loc='upper right', fontsize=10)
axes[1].grid(color='gray', linestyle='--', linewidth=0.5)

# Tight layout for better spacing
plt.tight_layout()

# Show the plots
plt.show()

```



```

In [23]: from sklearn.metrics import classification_report

# Generate predictions on the test set
y_pred = model.predict(test_generator)
y_pred_classes = np.argmax(y_pred, axis=1) # Convert probabilities to class indices
y_true = test_generator.classes # True class labels

# Class labels (assuming the generator's classes are in order of the directories)
class_labels = list(test_generator.class_indices.keys())

# Create the classification report
report = classification_report(y_true, y_pred_classes, target_names=class_labels)

# Print the report
print("Classification Report:")
print(report)

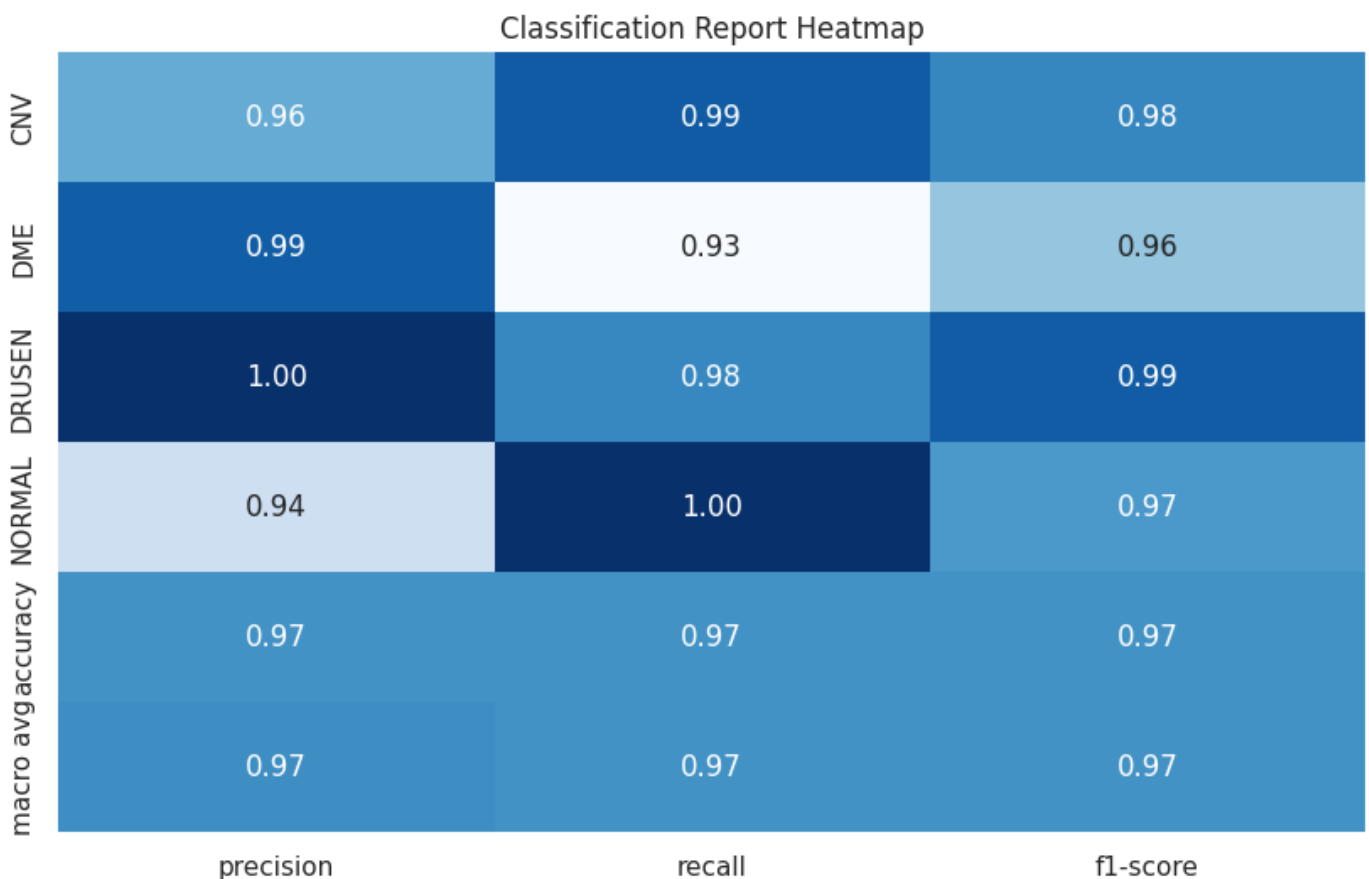
```

	precision	recall	f1-score	support
CNV	0.96	0.99	0.98	242
DME	0.99	0.93	0.96	242
DRUSEN	1.00	0.98	0.99	242
NORMAL	0.94	1.00	0.97	242
accuracy			0.97	968
macro avg	0.97	0.97	0.97	968
weighted avg	0.97	0.97	0.97	968

```
In [24]: import pandas as pd

# Convert the classification report to a dataframe
report_dict = classification_report(y_true, y_pred_classes, target_names=class_labels, o
report_df = pd.DataFrame(report_dict).transpose()


# Plot the heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(report_df.iloc[:-1, :-1], annot=True, fmt=".2f", cmap="Blues", cbar=False)
plt.title("Classification Report Heatmap")
plt.show()
```



```
In [27]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Generate confusion matrix
cm = confusion_matrix(y_true, y_pred_classes)

# Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_labels)
disp.plot(cmap='Blues', xticks_rotation=45)
plt.title("Confusion Matrix")
plt.grid(False)
plt.show()
```




```
In [37]: # Identify misclassified indices
misclassified_indices = np.where(y_true != y_pred_classes)[0]

# Plot misclassified images
plt.figure(figsize=(12, 12))
for i, idx in enumerate(misclassified_indices[:15]): # First 16 misclassified images
    plt.subplot(4, 4, i+1)
    plt.imshow(X_test[idx])
    plt.title(f"True: {class_labels[y_true[idx]]}, Pred: {class_labels[y_pred_classes[idx]]}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```

```
-----
IndexError                                Traceback (most recent call last)
Cell In[37], line 8
      6 for i, idx in enumerate(misclassified_indices[:15]): # First 16 misclassified i
      7     plt.subplot(4, 4, i+1)
----> 8     plt.imshow(X_test[idx])
      9     plt.title(f"True: {class_labels[y_true[idx]]}, Pred: {class_labels[y_pred_cl
asses[idx]]}")
     10     plt.axis('off')
```

IndexError: index 287 is out of bounds for axis 0 with size 194




```
In [34]: from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

# Binarize the labels for multi-class ROC
y_true_bin = label_binarize(y_true, classes=np.arange(len(class_labels)))

# Compute ROC curve and AUC for each class
fpr = {}
tpr = {}
roc_auc = {}
for i in range(len(class_labels)):
    fpr[i], tpr[i], _ = roc_curve(y_true_bin[:, i], y_pred[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot all ROC curves
plt.figure(figsize=(8, 6))
for i, label in enumerate(class_labels):
    plt.plot(fpr[i], tpr[i], label=f'{label} (AUC = {roc_auc[i]:.2f})')
plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.show()
```

```
In [14]: import tensorflow as tf
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import Dense, Flatten, GlobalAveragePooling2D, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.models import load_model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
In [15]: base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(128, 128, 3))

# Freeze the base model layers to retain pre-trained weights
base_model.trainable = True
```

```
In [16]: x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
output = Dense(4, activation='softmax')(x)

# Final model
model_resnet = Model(inputs=base_model.input, outputs=output)

# Compile the model
model_resnet.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

```
In [16]: history = model_resnet.fit(
    train_generator,
    validation_data=val_generator,
    epochs=10,
```

```
        verbose=1  
)
```

Epoch 1/10

2609/2609 ————— **9054s** 3s/step - accuracy: 0.8364 - loss: 0.5201 - val_accuracy: 0.9062 - val_loss: 0.3096

Epoch 2/10

2609/2609 ————— **8016s** 3s/step - accuracy: 0.9078 - loss: 0.2855 - val_accuracy: 1.0000 - val_loss: 0.0427

Epoch 3/10

2609/2609 ————— **8003s** 3s/step - accuracy: 0.9083 - loss: 0.2776 - val_accuracy: 0.9688 - val_loss: 0.0509

Epoch 4/10

2609/2609 ————— **8006s** 3s/step - accuracy: 0.9263 - loss: 0.2262 - val_accuracy: 0.9375 - val_loss: 0.1253

Epoch 5/10

2609/2609 ————— **7849s** 3s/step - accuracy: 0.9214 - loss: 0.2437 - val_accuracy: 1.0000 - val_loss: 0.0874

Epoch 6/10

2609/2609 ————— **7977s** 3s/step - accuracy: 0.9277 - loss: 0.2170 - val_accuracy: 1.0000 - val_loss: 0.0111

Epoch 7/10

2609/2609 ————— **7809s** 3s/step - accuracy: 0.9289 - loss: 0.2215 - val_accuracy: 1.0000 - val_loss: 0.0159

Epoch 8/10

2609/2609 ————— **8031s** 3s/step - accuracy: 0.9351 - loss: 0.1952 - val_accuracy: 1.0000 - val_loss: 0.0150

Epoch 9/10

2609/2609 ————— **7957s** 3s/step - accuracy: 0.9286 - loss: 0.2204 - val_accuracy: 1.0000 - val_loss: 0.0303

Epoch 10/10

2609/2609 ————— **7757s** 3s/step - accuracy: 0.9376 - loss: 0.1862 - val_accuracy: 1.0000 - val_loss: 0.0115

In [17]:

```
test_loss, test_accuracy = model_resnet.evaluate(test_generator)  
print(f"Test Loss: {test_loss}, Test Accuracy: {test_accuracy}")
```

31/31 ————— **30s** 978ms/step - accuracy: 0.9934 - loss: 0.0330

Test Loss: 0.031029945239424706, Test Accuracy: 0.9927685856819153